

Introduction to Performance Engineering

Nick Cameron



Disclaimers

- I'm not an expert
- Specific to general-purpose hardware
- Doesn't cover domain-specific issues




Background



Throughput

- Work done in a given time
- QPS




Latency

- Time taken by one unit of work
- μs



Macro-optimisations

- Potentially large gains
- Often trade-offs
- High risk



Micro-optimisations

- No trade-off
- Predictable
- Incremental
- Small gains
 - These compound:
 - 1% per week = 66% per year
 - (opposite is also true: -1%/week = -40%/year)

Hot code

- Focus on hot code (bottlenecks)
- Is it really hot?



Graphics example

60 fps

```
loop {  
    once_per_frame();  
    ...  
}
```



Graphics example

60 fps

16.7 ms

~2 million instructions

```
loop {  
    once_per_frame();  
    ...  
}
```



Graphics example

4k, 60 fps

```
loop {  
  for x in 0..width {  
    for y in 0..height {  
      once_per_pixel();  
    }  
  }  
  ...  
}
```



Graphics example

4k, 60 fps

0.5bn pixels per second

1.9ns

60 instructions

```
loop {  
  for x in 0..width {  
    for y in 0..height {  
      once_per_pixel();  
    }  
  }  
  ...  
}
```



Big O

- $O(n)$, $O(e^n)$, etc
- Growth which dominates for large n



Big O

- How big is your n ?
 - $O(1)$
 - $O(n)$, $O(n^2)$, ...
 - $O(e^n)$



Big O

- Worst case
 - average case?
 - amortised cost?
- How big is the constant factor?



Computer architecture






```
lea ecx, [rsi - 1]  
lea edx, [rsi - 2]  
imul rdx, rcx  
shr rdx
```





42













lea ecx, [rsi - 1]

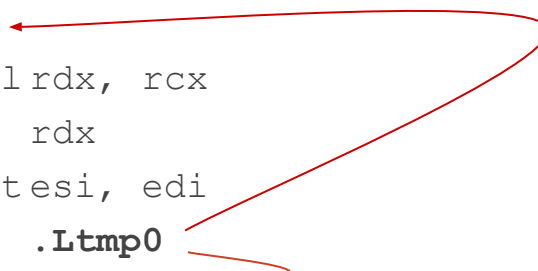









Branch prediction


```
    lea edx, [rsi - 2]
.Ltmp0:
    imul rdx, rcx
    shr rdx
    test esi, edi
    je .Ltmp0
    add eax, ecx
```












Cache rules everything around me

- Memory is the new IO
- Cache behaviour is the most important aspect of performance





Cache line width



Cache line width




Cache line width




Cache line width

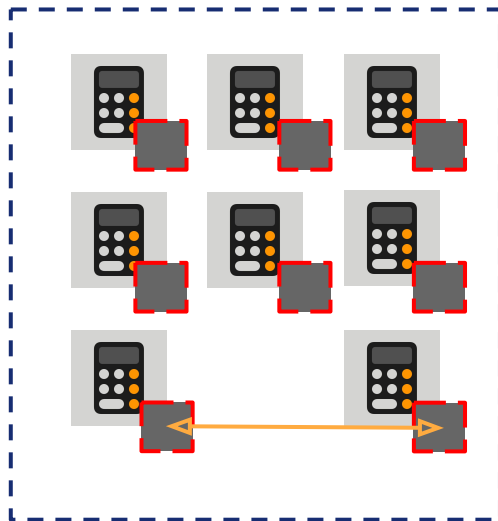


Cache line width



Cache line width





Optimising compilers

- `x * 4 => x << 2`
- `for i in 0..8 { x += i * MAGIC_NUMBER; } => 140`



Optimising compilers

- Inlining
- Compilation units



Benchmarking and profiling



Benchmarks

- You have to measure



Benchmarks

- You have to measure!



Benchmarks

- Reproducible
- Realistic



Benchmarks

- How much of the system are you simulating?
- What is your input data?
- Statistical issues and warm-up
- Environment
- Are you optimizing for the benchmark?



Profiling

- perf
- valgrind
- vtune
- flamegraphs
- metrics
- instrumentation



Profiling

- Indirect costs
- Wait time




Optimisation



The real problem is that programmers have spent far too much time worrying about efficiency in the wrong places and at the wrong times; **premature optimization is the root of all evil** (or at least most of it) in programming.

Donald Knuth



Premature optimisation

- Costs to optimisation
- You may be wrong
- Simple, clean code is easier to optimise
- CPUs and compilers are pretty smart



Premature optimisation

- Simple, clear code
- Don't write sloppy code
 - Think about algorithms and concurrency
 - Avoid unnecessary costs
 - Don't hide costs



Optimisation

- Measure
- Find hot code
 - Make it faster
 - Make it less hot
- Measure again



Optimisation

- Measure
- Find hot code
 - Make it faster
 - **Make it less hot**
- Measure again



Caching

- Memoization
- Dynamic programming




Batching



- Many operations
- Minimise overhead



Batching



Batching



Optimisation


- Measure
- Find hot code
 - Make it faster
 - Make it less hot
- Measure again





IO

- No.






Concurrency



Concurrency

- No sharing is best.
- Immutable sharing is fine.
 - Split mutable state
- Atomics.
- Locks, etc.





Cache-friendly programming

- Sequential data.
- Small, tightly packed.
- Predictable code.
- **Indirection.**



Cache-friendly programming

- Arrays, slices, `Vec`
- **Linked lists, graphs**



Cache-friendly programming

- Data layout
- Data-oriented design



Cache-friendly programming

- Data layout
- Data-oriented design


```
struct Foo {  
    x: u32,  
    y: u32,  
    val: u64,  
}
```



Cache-friendly programming

- Data layout
- Data-oriented design





```
struct Foo {  
    x: u32,  
    y: u32,  
    val: u64,  
}
```



Cache-friendly programming

- Data layout
- Data-oriented design

```
struct Foo {  
    x: u32,  
    y: u32,  
    val: u64,  
}
```



Allocation

- Heap vs stack
 - Dynamic vs static
 - Bookkeeping
 - Free
 - GC
- Cache effects
 - size trade-off



Predictable code

- Few branches per loop
- Branch the same way
- Pre-sort mixed data



Predictable code

```
for x in ... {  
  if x.abc < 0 {  
    ...  
  }  
  if foo(x.xyz) {  
    ...  
  }  
}
```



Predictable code

```
for x in ... {  
    if x.abc < 0 {  
        ...  
    }  
}  
...  
for x in ... {  
    if foo(x.xyz) {  
        ...  
    }  
}
```



Tuning

- OS
- compiler
- runtime
- allocator
- parameters



Summary

- **You have to measure**
- Keep it simple
- Parallelism
- Cache effects
- Allocation



Performant Rust



Concurrency

- TLS
- Crossbeam



Concurrency

- async is fast, but introduces complexity
- worthwhile if you have *lots* of concurrency and lots of time spent blocking on io



Allocation

- `Box, Vec, String, ...`
- `collect, to_string, ...`



Allocation

- Trade-off between size and indirection
- Data structures
 - hybrid data structures
 - `with_capacity`
- Arena allocation



Large data

- enum variants
- `repr("C")`
- `-Zprint-type-sizes`



Dispatch

- `fn foo(x: &impl Bar)`
- `fn foo<X: Bar>(x: &X)`



Dispatch

- `fn foo(x: &impl Bar)`
- `fn foo<X: Bar>(x: &X)`
- `fn foo(x: &dyn Bar)`
- `fn foo(x: &Bar)`



Dispatch

- `fn foo(x: &impl Bar)`
 - thin pointer
 - static dispatch
 - monomorphised



Dispatch

- `fn foo(x: &dyn Bar)`
 - fat pointer
 - virtual dispatch
 - polymorphic generated code



Dispatch

- monomorphised code
 - code bloat
 - how big is the function?
 - how many types is it used with?



Dispatch

- virtual dispatch
 - indirect call
 - is it predictable?



Data types

- There are many variations
 - different structures/algorithms
 - different implementations
 - different constraints
- Don't pay for what you don't need
 - Cryptographic hash
 - Secure random numbers
 - ...



Run time and compile time

- Bounds checks are expensive
 - use iteration
 - if you can't iterate, assert bounds up front



Run time and compile time

- Iterators
 - `chain`
 - `retain, partition, ...`
 - `collect`



Run time and compile time

- `Arc` and `Rc`
- `Mutex` and `RefCell`
- overflow checks



Run time and compile time

- Compile time
 - is free! (-ish)
 - const fn
 - procedural macros, build scripts
- `assert` and `debug_assert`



Run time and compile time

- Compile time
 - compile times
 - code bloat
 - macro-generated code



The compiler, etc

- rustc and llvm
- Inlining
 - `#[inline]`
 - recursion
- Panicking
 - overheads
 - panic = abort



The compiler, etc

- SIMD



The compiler, etc

- LTO
- # [bench], Criterion
- -Z



Thank You !
谢谢



Remember to measure 🤔

