TiDB TechDay 2019

The Evolution of TiKV

How can we achieve huge performance improvements?

Agenda

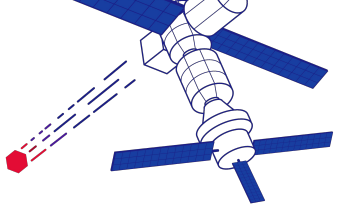

1. **Achievement**
2. Future Plan
3. Community

Performance Problems

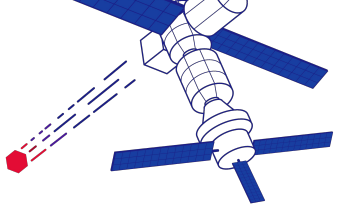

1. Single thread
2. RocksDB write amplification
3. Too many Region heartbeats
4. GC lagging
5. Big query
6. Rough scheduling
7. gRPC
8. Em…

Threaded Raft

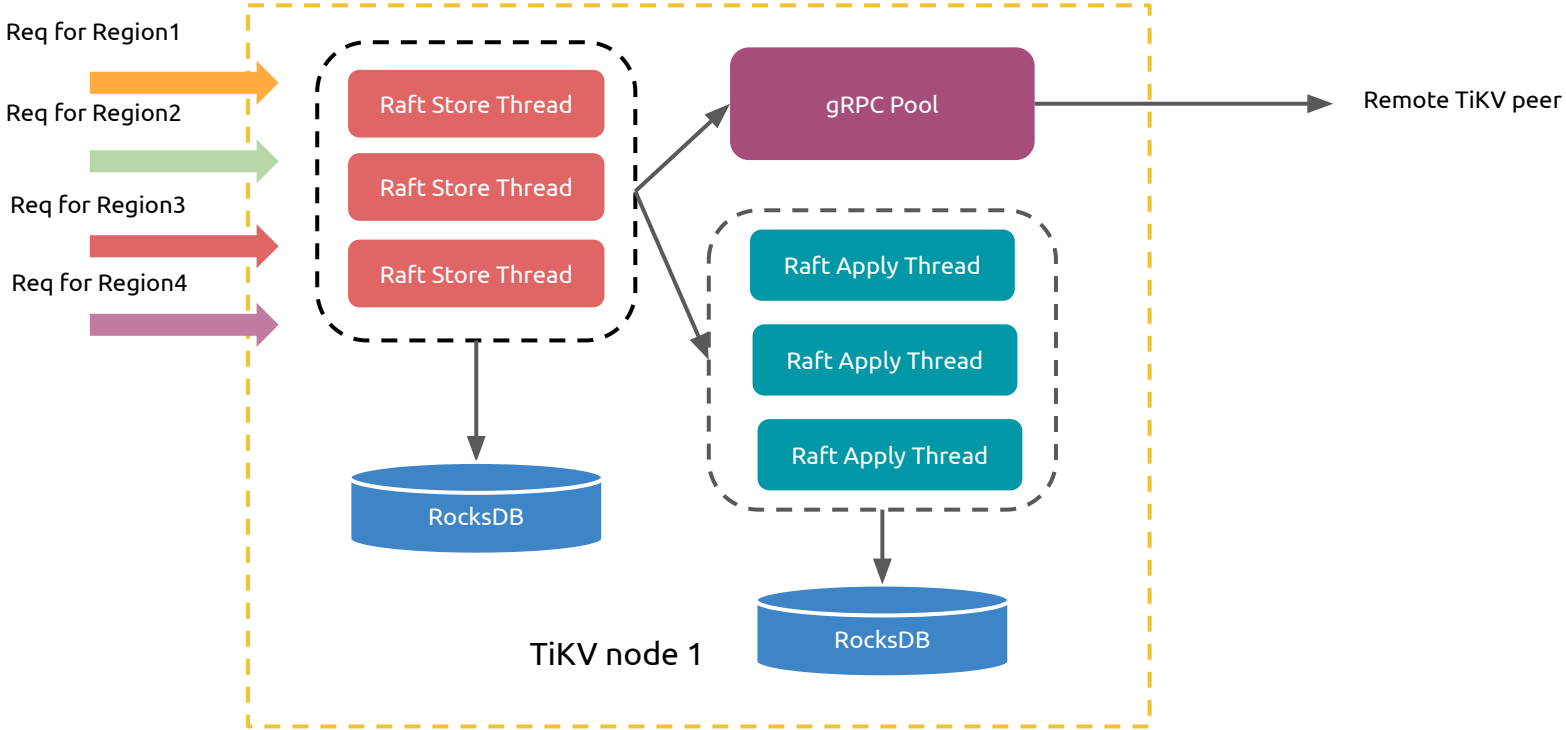

Hibernate Region

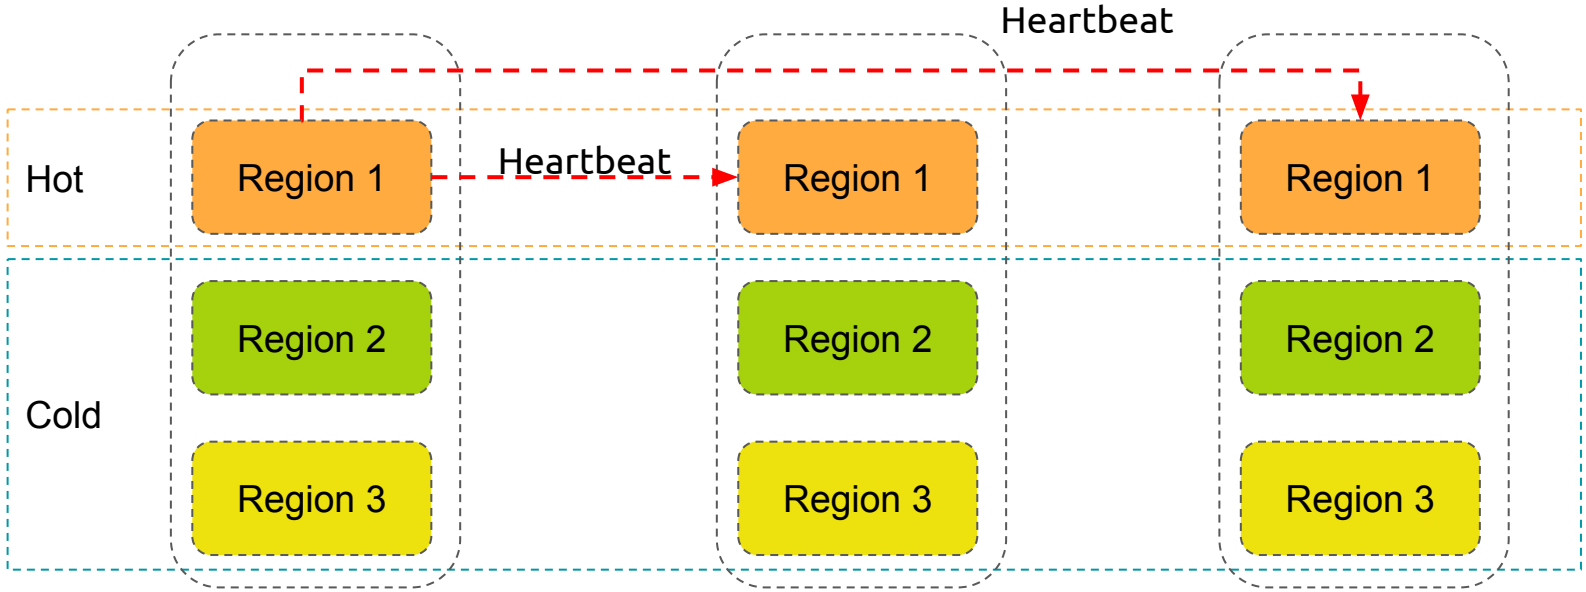

Distributed GC

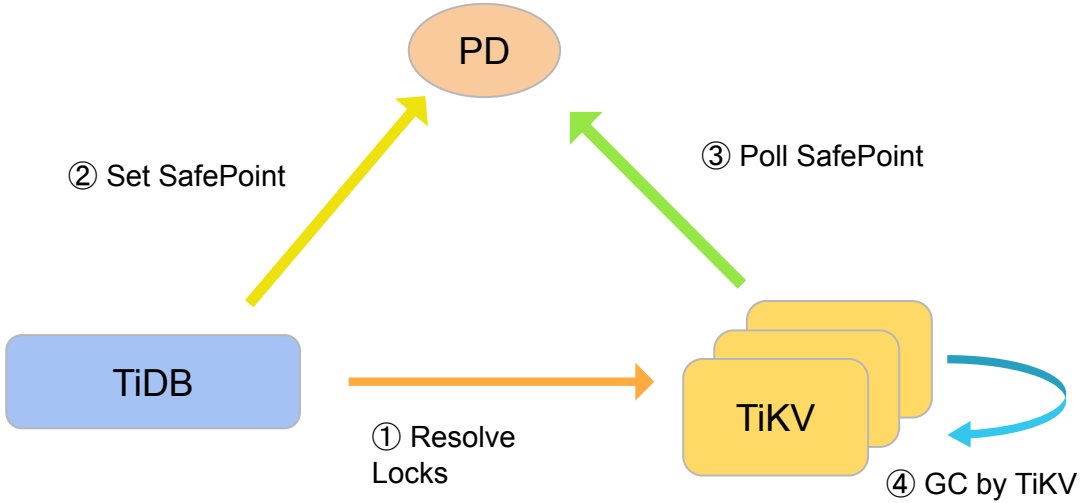

Coprocessor Vectorization

```
SELECT SUM(foo) FROM Table WHERE foo>bar
```

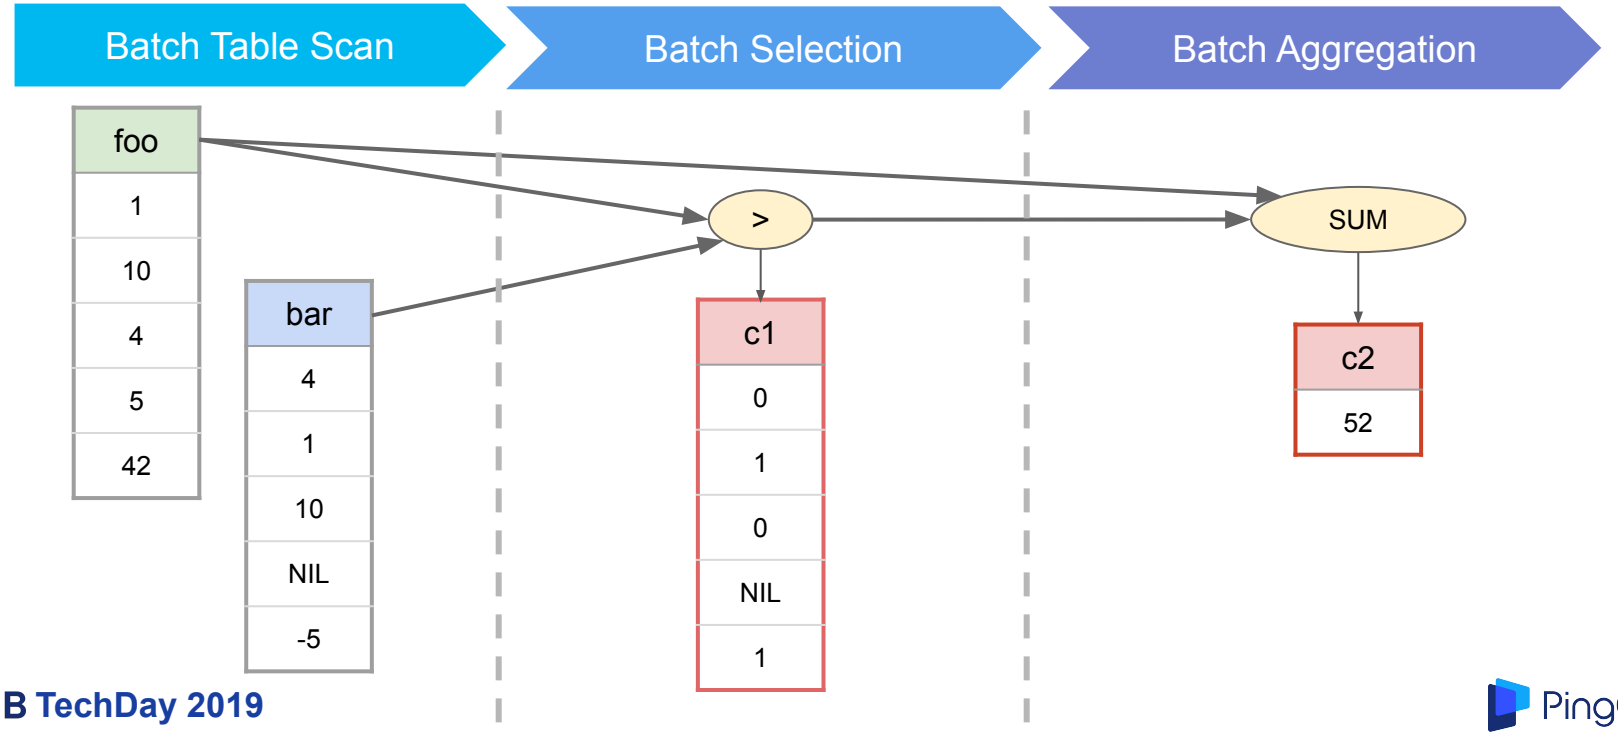

Store Limit Scheduler

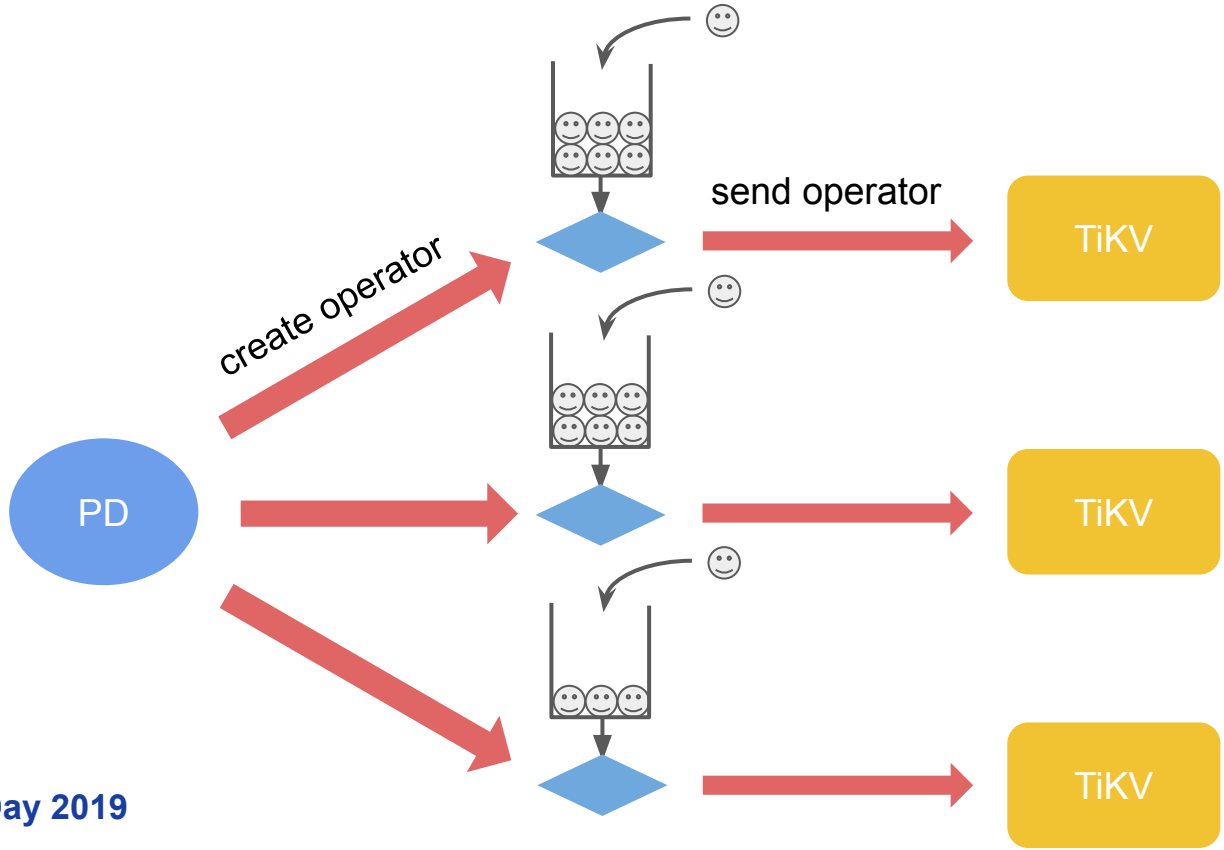

gRPC Batch messages

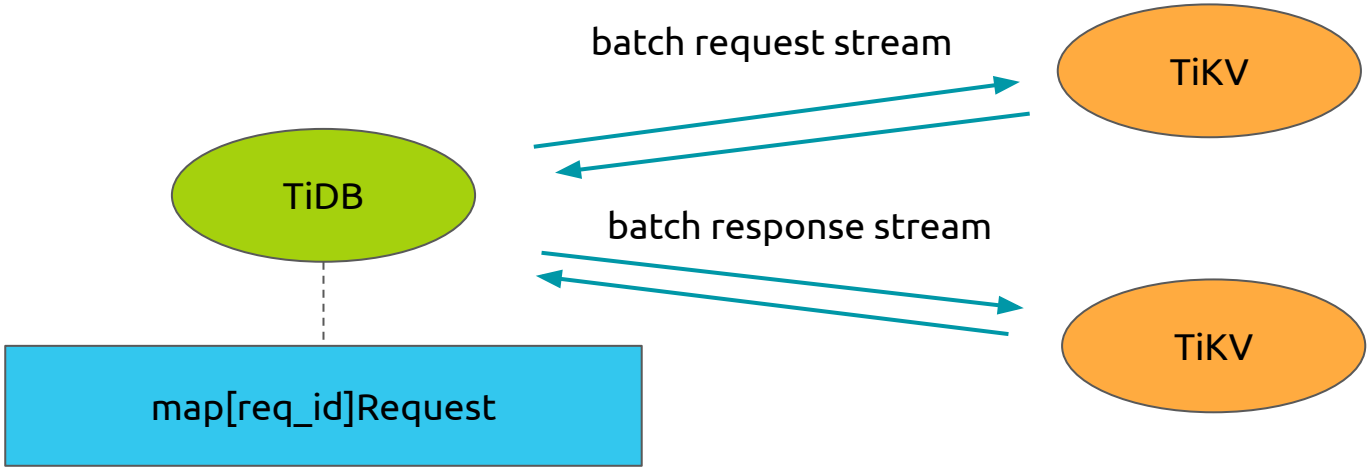

Titan

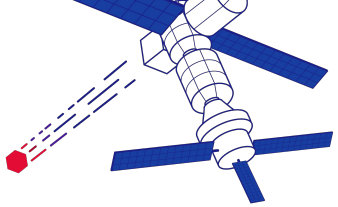

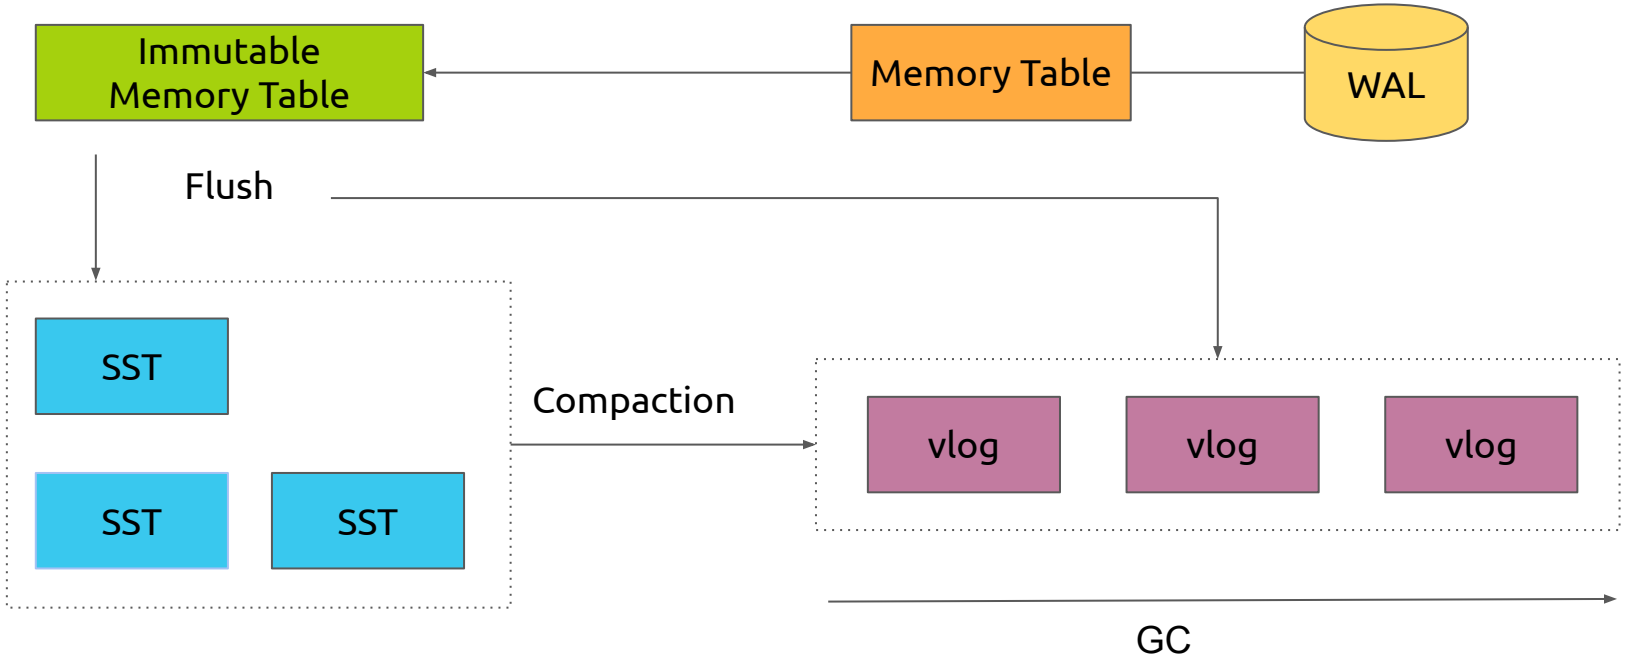
Immutable
Memory Table
Memory Table
WAL
Flush
SST
SST
SST
Compaction
vlog
vlog
vlog
GC

Improvements - Sysbench

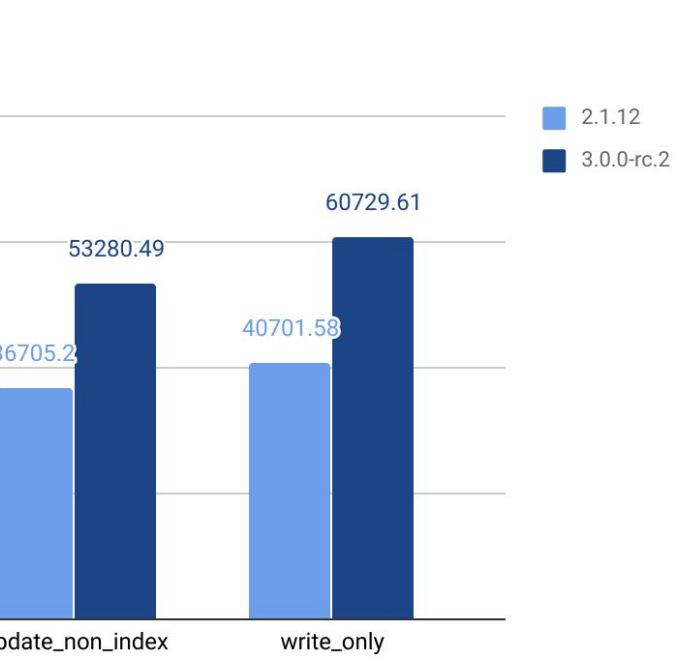

Improvements - YCSB

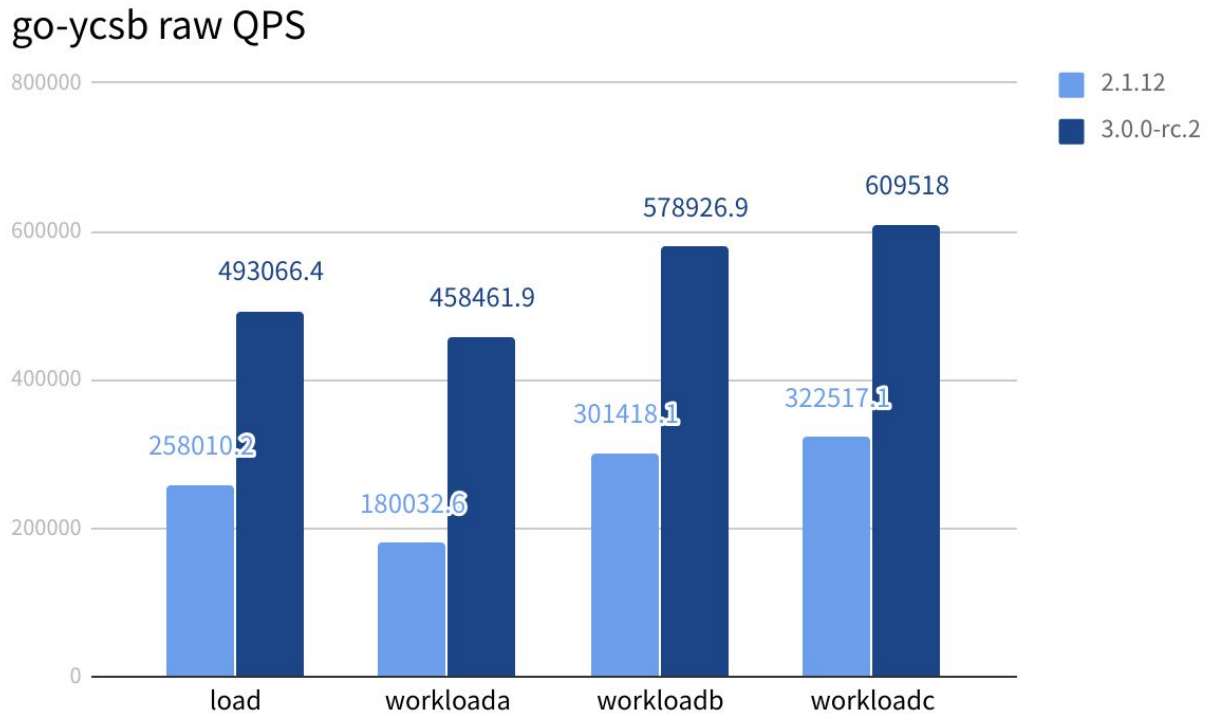

Improvements - TPC-H

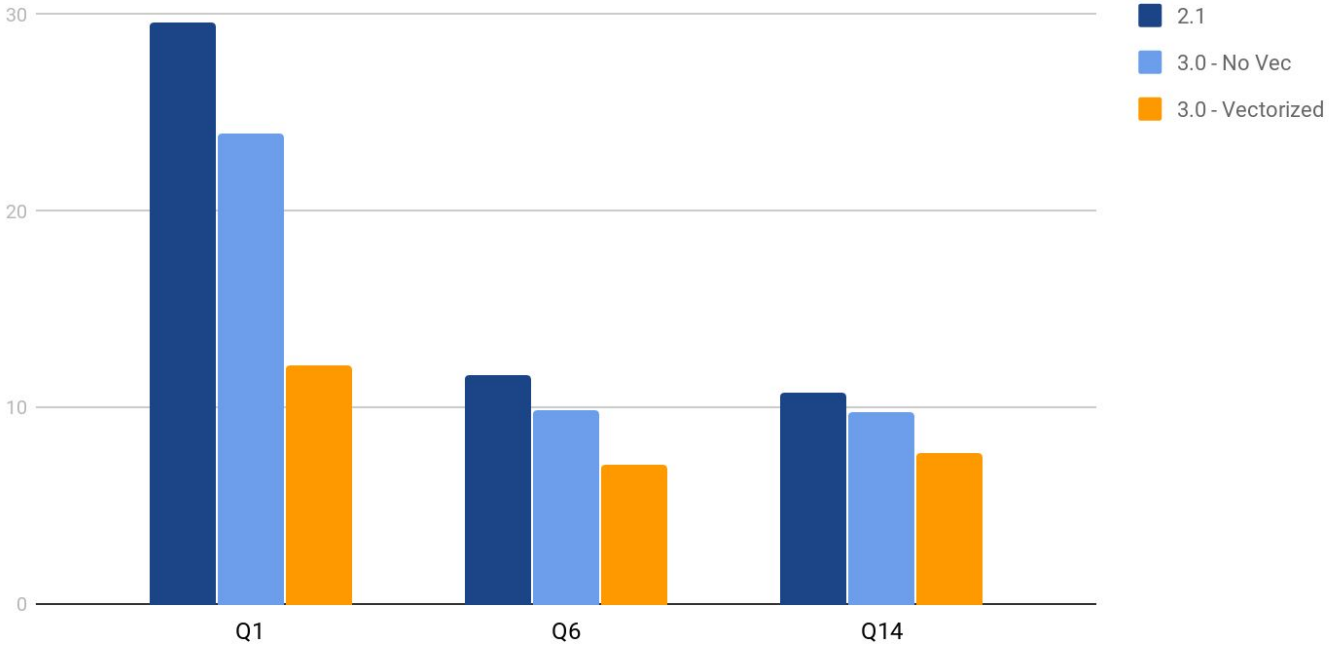

Improvements - Titan

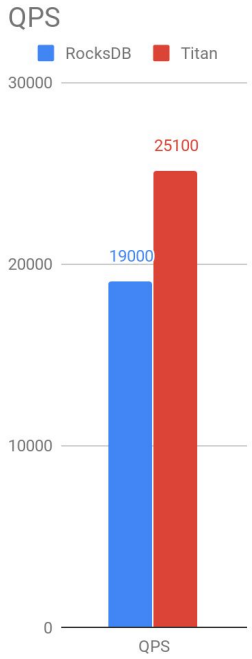

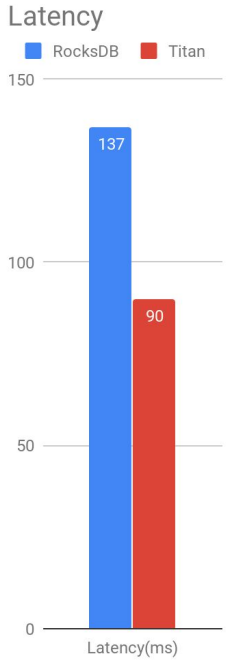

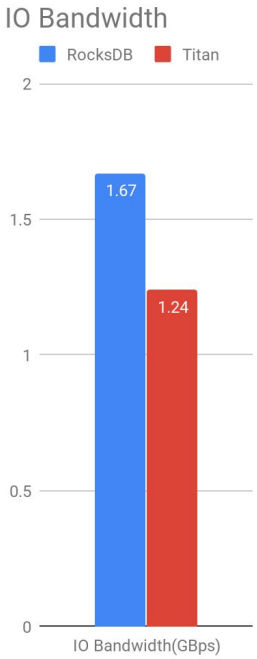

Improvements - Scheduling

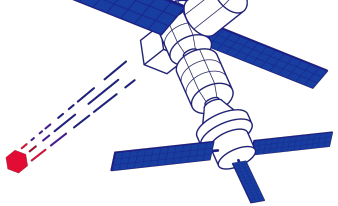

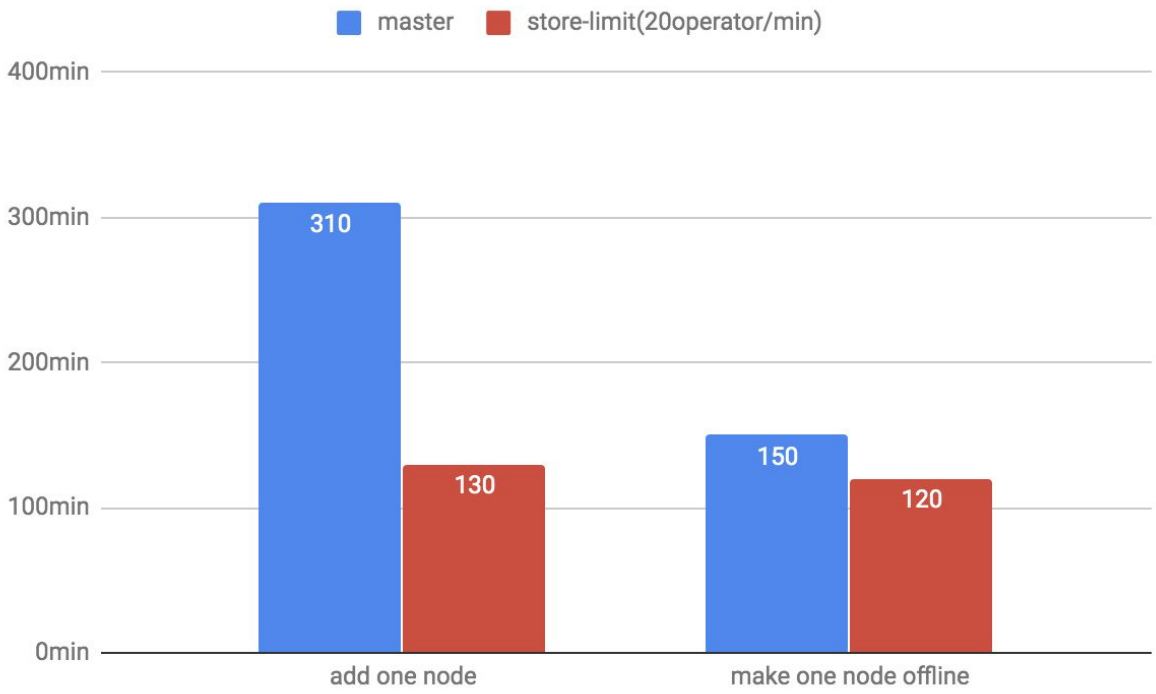

Agenda

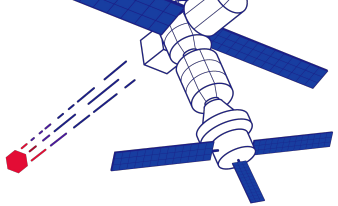

1. Achievement
2. **Future Plan**
3. Community

Raft - Chain Replication

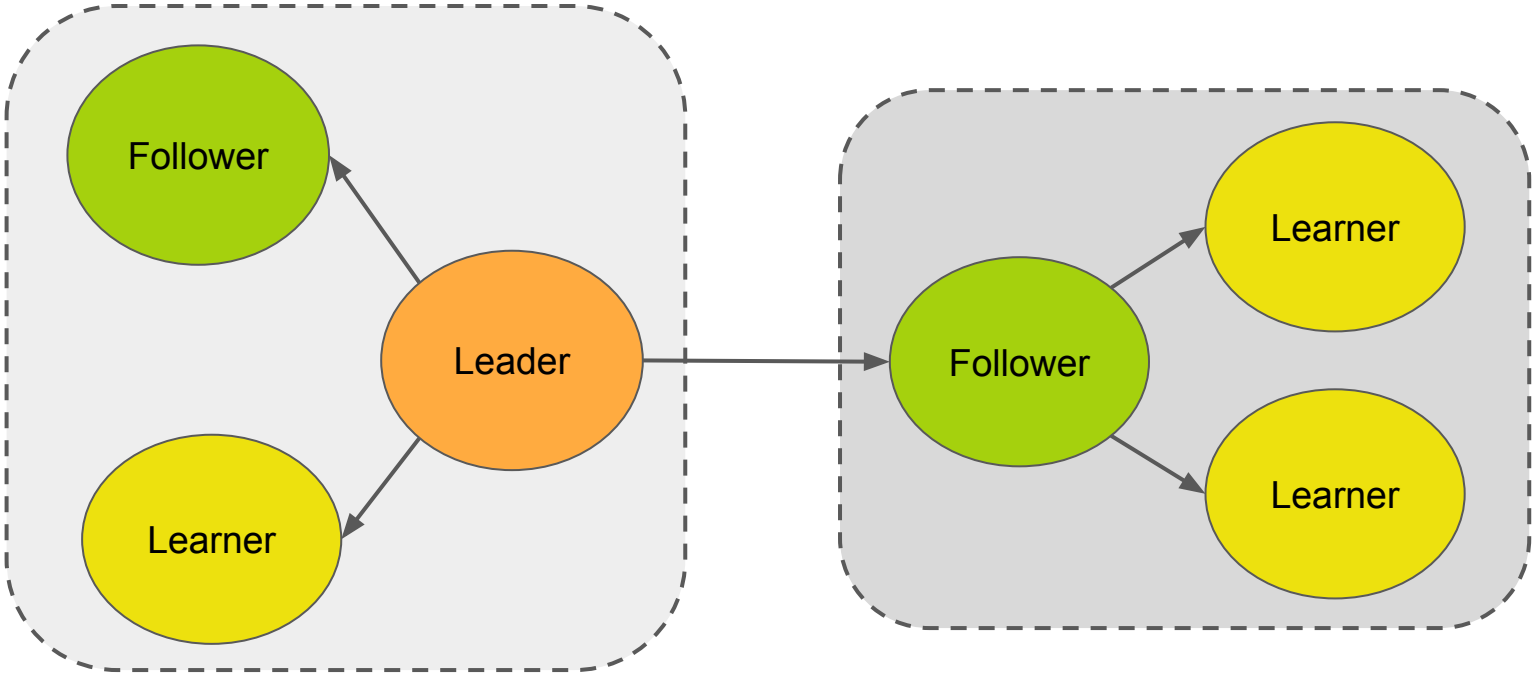

Transaction - 1 PC

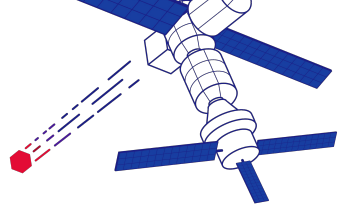

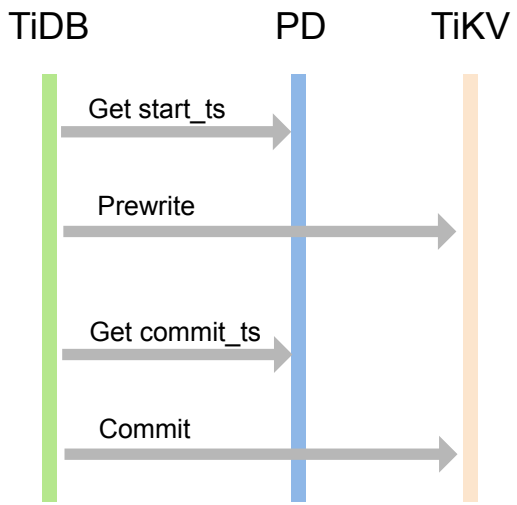

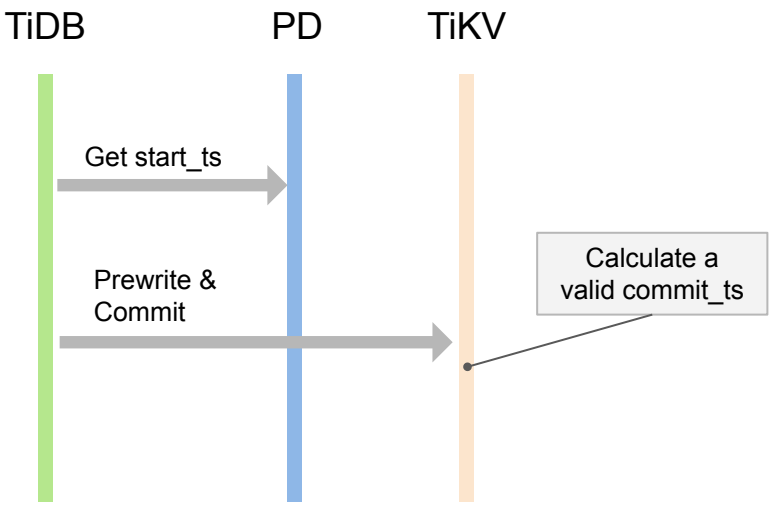

When there's only one Region affected by the transaction...

Backup, Restore and Replication

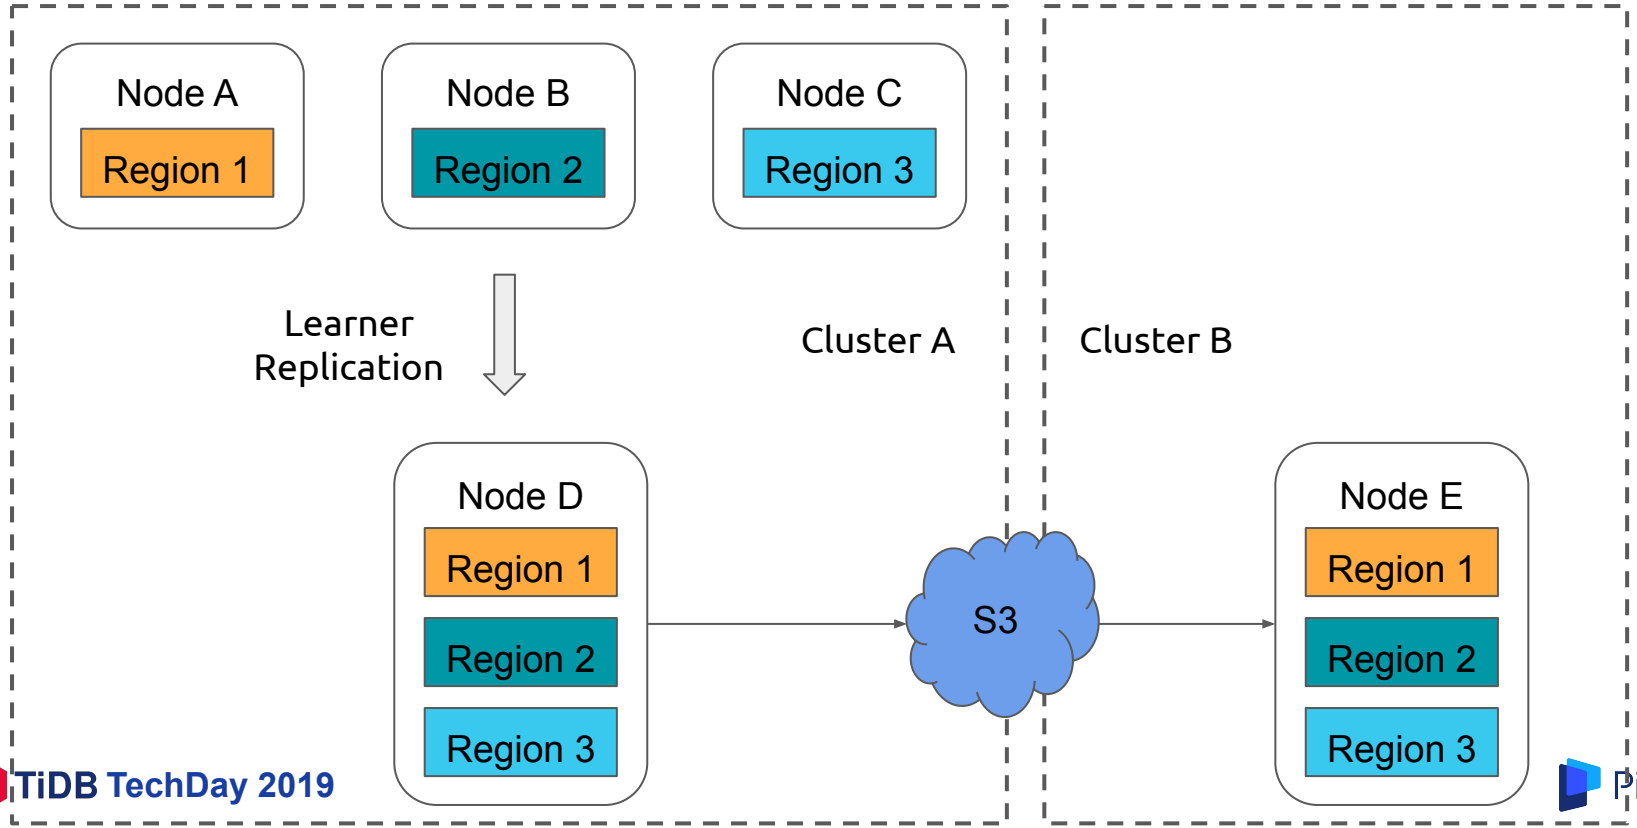

Elastic Scheduling

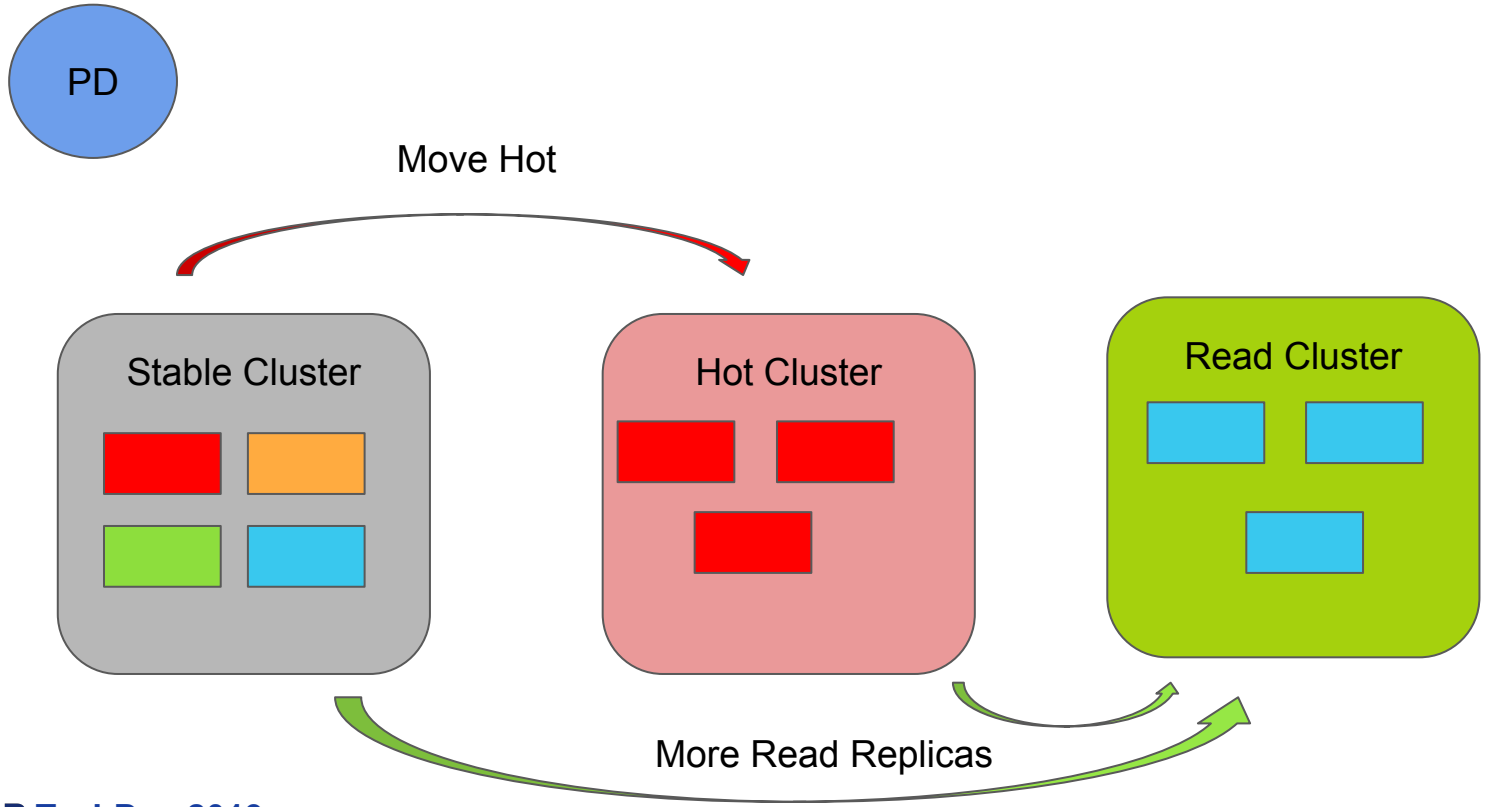

RocksDB Guard

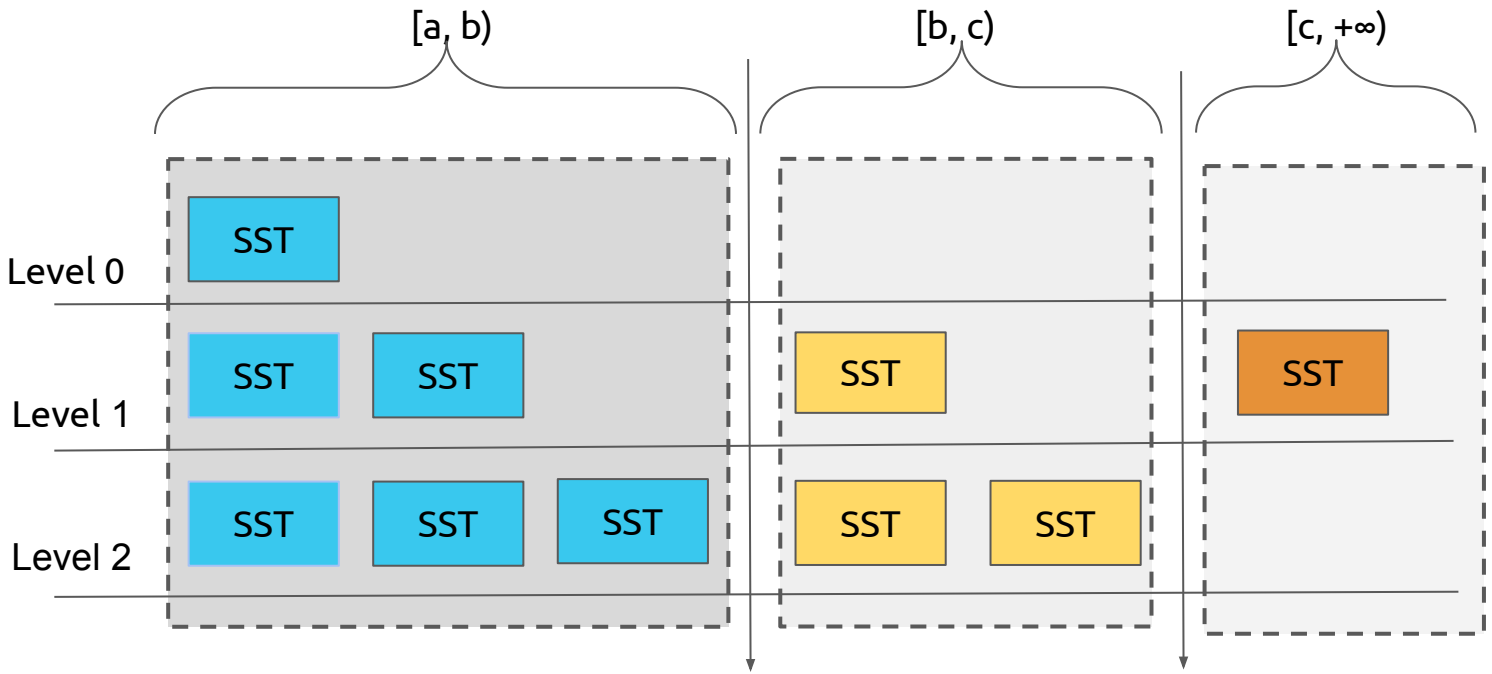

Agenda

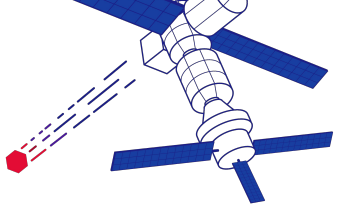

1. Achievement
2. Future Plan
3. **Community**

CNCF

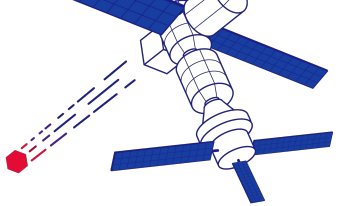

CLOUD NATIVE
COMPUTING FOUNDATION

TiKV

Timeline

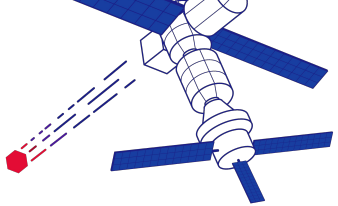

- July 2018: First presentation to CNCF Technical Oversight Committee (TOC)
- August 2018: Enters CNCF as a Sandbox level member project
- April 2019: Begins technical review and due diligence for Incubating status
- May 2019: TiKV reaches Incubating status

Talent Plan

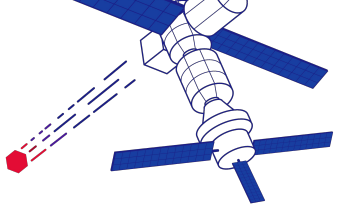

- Practical Networked Applications in Rust
- Distributed Systems in Rust

Test

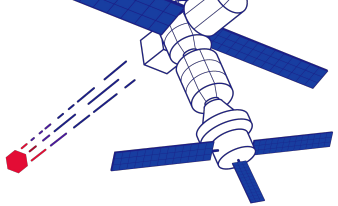

1. jepsen-io/jepsen
2. pingcap/chaos

Thank You !